

CS 109A/STAT 121A/AC 209A/CSCI E-109A: Homework 3

Multiple Linear Regression, Subset Selection, Cross Validation

Harvard University

Fall 2017

Instructors: Pavlos Protopapas, Kevin Rader, Rahul Dave, Margo Levine

INSTRUCTIONS

- To submit your assignment follow the instructions given in canvas.
 - Restart the kernel and run the whole notebook again before you submit.
 - Do not include your name(s) in the notebook if you are submitting as a group.
 - If you submit individually and you have worked with someone, please include the name of your [one] partner below.
-

Your partner's name (if you submit separately): Group 22

Enrollment Status (109A, 121A, 209A, or E109A): AC209A

Import libraries:

```
In [1]: import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
from sklearn.metrics import r2_score
import statsmodels.api as sm
from statsmodels.api import OLS
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.linear_model import RidgeCV
from sklearn.linear_model import LassoCV
%matplotlib inline

import seaborn as sns
```

```
D:\Anaconda3\envs\py36\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The pandas.core.date
tools module is deprecated and will be removed in a future version. Please use the pandas.tseries module inst
ead.
```

```
from pandas.core import datetools
```

Forecasting Bike Sharing Usage

In this homework, we will focus on multiple linear regression and will explore techniques for subset selection. The specific task is to build a regression model for a bike share system that can predict the total number of bike rentals in a given day, based on attributes about the day. Such a demand forecasting model would be useful in planning the number of bikes that need to be available in the system on any given day, and also in monitoring traffic in the city. The data for this problem was collected from the Capital Bikeshare program in Washington D.C. over two years.

The data set is provided in the files `Bikeshare_train.csv` and `Bikeshare_test.csv`, as separate training and test sets. Each row in these files contains 10 attributes describing a day and its weather:

- season (1 = spring, 2 = summer, 3 = fall, 4 = winter)
- month (1 through 12, with 1 denoting Jan)
- holiday (1 = the day is a holiday, 0 = otherwise)
- day_of_week (0 through 6, with 0 denoting Sunday)
- workingday (1 = the day is neither a holiday or weekend, 0 = otherwise)
- weather
 - 1: Clear, Few clouds, Partly cloudy, Partly cloudy
 - 2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist
 - 3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds
 - 4: Heavy Rain + Ice Pallets + Thunderstorm + Mist, Snow + Fog
- temp (temperature in Celsius)
- atemp (apparent temperature, or relative outdoor temperature, in Celsius)
- humidity (relative humidity)
- windspeed (wind speed)

and the last column 'count' contains the response variable, i.e. total number of bike rentals on the day.

Part (a): Data Exploration & Preprocessing

As a first step, identify important characteristics of the data using suitable visualizations when necessary. Some of the questions you may ask include (but are not limited to):

- How does the number of bike rentals vary between weekdays and weekends?
- How about bike rentals on holidays?
- What effect does the season have on the bike rentals on a given day?
- Is the number of bike rentals lower than average when there is rain or snow?
- How does temperature effect bike rentals?
- Do any of the numeric attributes have a clear non-linear dependence with number of the bike rentals?

```
In [2]: dtype_dict = {"season":int, "month":int, "holiday":int, "day_of_week": int, "workingday":int, "weather":int, \
                    "temp":float, "atemp":float, "humidity":float, "windspeed": float, "count":int}
```

```
In [3]: bike_train_df = pd.read_csv('data/Bikeshare_train.csv', low_memory=False, index_col=0, dtype=dtype_dict)
bike_test_df = pd.read_csv('data/Bikeshare_test.csv', low_memory=False, index_col=0, dtype=dtype_dict)
bike_train_df.head()
```


Out[3]:

	season	month	holiday	day_of_week	workingday	weather	temp	atemp	humidity	windspeed	count
0	2	5	0	2	1	2	24.0	26.0	76.5833	0.118167	6073
1	4	12	0	2	1	1	15.0	19.0	73.3750	0.174129	6606
2	2	6	0	4	1	1	26.0	28.0	56.9583	0.253733	7363
3	4	12	0	0	0	1	0.0	4.0	58.6250	0.169779	2431
4	3	9	0	3	1	3	23.0	23.0	91.7083	0.097021	1996

```
In [4]: bike_count_train = bike_train_df['count'].values
bike_count_test = bike_test_df['count'].values
```

```
In [5]: season = bike_train_df['season'].values

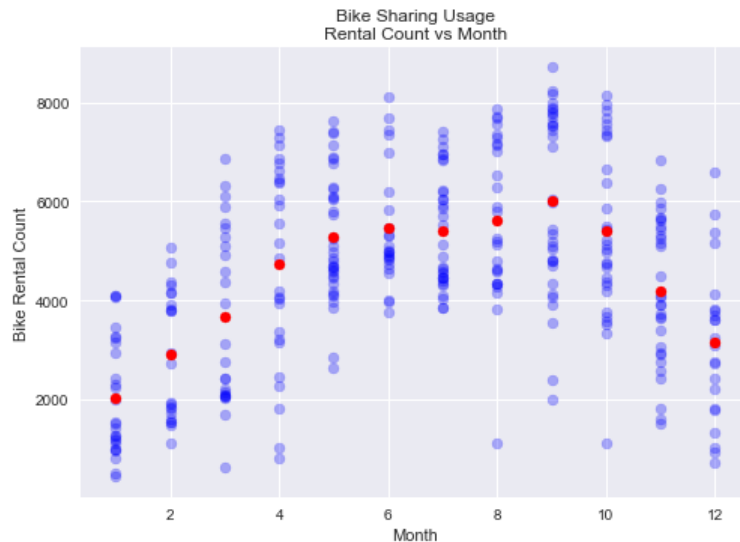
sns.set_context('notebook')
plt.scatter(season, bike_count_train, color='blue', alpha=0.3, label='data')
for s in range(1, 5):
    plt.scatter(s, bike_count_train[season==s].mean(), color='red', label='mean')
plt.xlabel('Season')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Season')
plt.show()
```



The number of bike rentals seems to have some relationship with the season. Rental count is higher in Fall (s=3) than in Spring (s=1)

```
In [6]: month = bike_train_df['month'].values

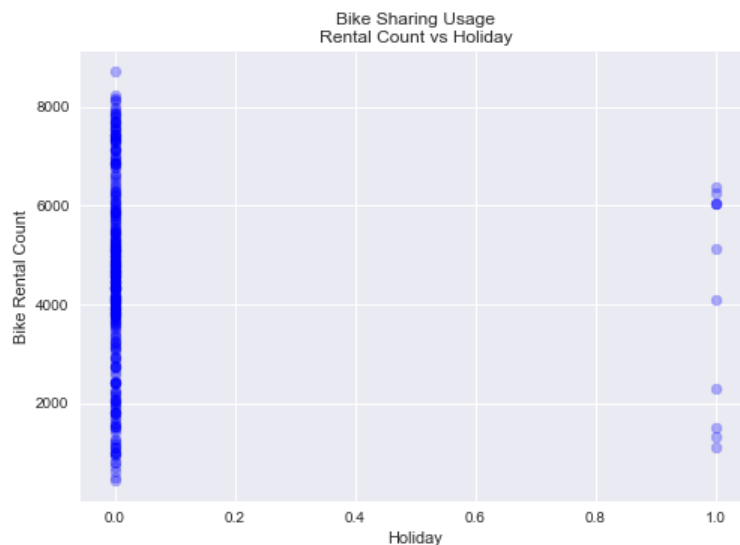
sns.set_context('notebook')
plt.scatter(month, bike_count_train, color='blue', alpha=0.3)
for s in range(1, 13):
    plt.scatter(s, bike_count_train[month==s].mean(), color='red', label='mean')
plt.xlabel('Month')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Month')
plt.show()
```



The number of bike rentals is relatively higher during summer (s=5-10) and lower during the colder period

```
In [7]: holiday = bike_train_df['holiday'].values

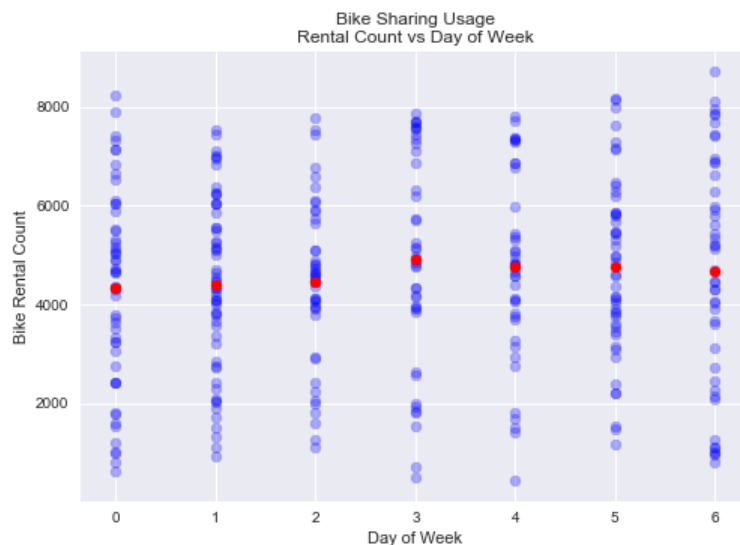
sns.set_context('notebook')
plt.scatter(holiday, bike_count_train, color='blue', alpha=0.3)
plt.xlabel('Holiday')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Holiday')
plt.show()
```



Too few data point in Holiday to make any meaningful conclusion

```
In [8]: day_of_week = bike_train_df['day_of_week'].values

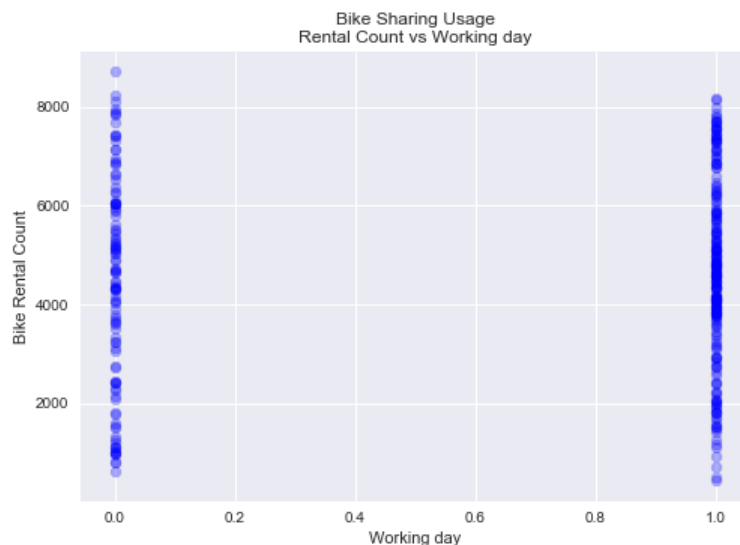
sns.set_context('notebook')
plt.scatter(day_of_week, bike_count_train, color='blue', alpha=0.3)
for s in range(0, 7):
    plt.scatter(s, bike_count_train[day_of_week==s].mean(), color='red', label='mean')
plt.xlabel('Day of Week')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Day of Week')
plt.show()
```



The day of week doesn't seem to affect the number of bike rental

```
In [9]: workingday = bike_train_df['workingday'].values

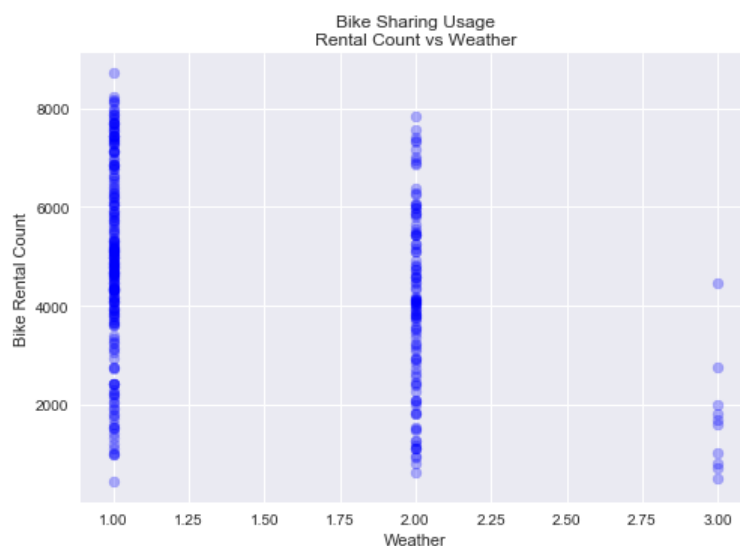
sns.set_context('notebook')
plt.scatter(workingday, bike_count_train, color='blue', alpha=0.3)
plt.xlabel('Working day')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Working day')
plt.show()
```



Whether it is a working day or not doesn't seem to affect the number of bike rental

```
In [10]: weather = bike_train_df['weather'].values


sns.set_context('notebook')
plt.scatter(weather, bike_count_train, color='blue', alpha=0.3)
plt.xlabel('Weather')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Weather')
plt.show()
```



During better weather (1 and 2) there is more rental counts, but there is too few data in bad weather (3 or 4) to conclude.

```
In [11]: temp = bike_train_df['temp'].values


sns.set_context('notebook')
plt.scatter(temp, bike_count_train, color='blue', alpha=0.5)
plt.xlabel('Temperature')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Temperature')
plt.show()
```



The rental count is relatively higher in higher temperature (20-30) and lower in low temperature (<10)

```
In [12]: atemp = bike_train_df['atemp'].values


sns.set_context('notebook')
plt.scatter(atemp, bike_count_train, color='blue', alpha=0.5)
plt.xlabel('Apparent Temperature')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Apparent Temperature')
plt.show()
```



The rental count is relatively higher in higher apparent temperature (20-30) and lower in low apparent temperature (<20)

```
In [13]: humidity = bike_train_df['humidity'].values


sns.set_context('notebook')
plt.scatter(humidity, bike_count_train, color='blue', alpha=0.5)
plt.xlabel('Humidity')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Humidity')
plt.show()
```



The scatter plot looks pretty random - most likely humidity has very small effect on rental count

```
In [14]: windspeed = bike_train_df['windspeed'].values

sns.set_context('notebook')
plt.scatter(windspeed, bike_count_train, color='blue', alpha=0.5)
plt.xlabel('Windspeed')
plt.ylabel('Bike Rental Count')
plt.title('Bike Sharing Usage\nRental Count vs Windspeed')
plt.show()
```



The scatter plot looks pretty random - most likely the windspeed has very small effect on rental count

We next require you to pre-process the categorical and numerical attributes in the data set:

- Notice that this data set contains categorical attributes with two or more categories. **Why can't they be directly used as predictors?** Convert these categorical attributes into multiple binary attributes using one-hot encoding: in the place of every categorical attribute x_j that has categories $1, \dots, K_j$, introduce $K_j - 1$ binary predictors $x_{j1}, \dots, x_{j,K_j-1}$ where x_{jk} is 1 whenever $x_j = k$ and 0 otherwise. **Why is it okay to not have a binary column for the K_j -th category?**
- Since the attributes are in different scales, it is a good practice to standardize the continuous predictors, i.e. to scale each continuous predictor to have zero mean and a standard deviation of 1. This can be done by applying the following transform to each continuous-valued predictor j : $\hat{x}_{ij} = (x_{ij} - \bar{x}_j)/s_j$, where \bar{x}_j and s_j are the sample mean and sample standard deviation (SD) of predictor j in the training set. We emphasize that the mean and SD values used for standardization must be estimated using only the training set observations, while the transform is applied to both the training and test sets. **Why shouldn't we include the test set observations in computing the mean and SD?**
- Provide a table of the summary statistics of the new attributes ('pd.describe()' function will help).

Hint: You may use the `pd.get_dummies` function to convert a categorical attribute in a data frame to one-hot encoding. This function creates K binary columns for an attribute with K categories. We suggest that you delete the last (or first) binary column generated by this function.

Note: We shall use the term "attribute" to refer to a categorical column in the data set, and the term "predictor" to refer to the individual binary columns resulting out of one-hot encoding.

```
In [15]: bike_train_df = pd.get_dummies(bike_train_df, columns = ['season', 'month', 'day_of_week', 'weather'])
bike_test_df = pd.get_dummies(bike_test_df, columns = ['season', 'month', 'day_of_week', 'weather'])
bike_train_df.head()
```

Out[15]:

	holiday	workingday	temp	atemp	humidity	windspeed	count	season_1	season_2	season_3	...	day_of_week_0	day_c
0	0	1	24.0	26.0	76.5833	0.118167	6073	0	1	0	...	0	0
1	0	1	15.0	19.0	73.3750	0.174129	6606	0	0	0	...	0	0
2	0	1	26.0	28.0	56.9583	0.253733	7363	0	1	0	...	0	0
3	0	0	0.0	4.0	58.6250	0.169779	2431	0	0	0	...	1	0
4	0	1	23.0	23.0	91.7083	0.097021	1996	0	0	1	...	0	0

5 rows × 33 columns

```
In [16]: # Drop redundant columns
bike_train_df = bike_train_df.drop(['season_4', 'month_12', 'day_of_week_0'], axis=1)
bike_test_df = bike_test_df.drop(['season_4', 'month_12', 'day_of_week_0'], axis=1)
bike_train_df.head()
```

Out[16]:

	holiday	workingday	temp	atemp	humidity	windspeed	count	season_1	season_2	season_3	...	month_11	day_of_wec
0	0	1	24.0	26.0	76.5833	0.118167	6073	0	1	0	...	0	0
1	0	1	15.0	19.0	73.3750	0.174129	6606	0	0	0	...	0	0
2	0	1	26.0	28.0	56.9583	0.253733	7363	0	1	0	...	0	0
3	0	0	0.0	4.0	58.6250	0.169779	2431	0	0	0	...	0	0
4	0	1	23.0	23.0	91.7083	0.097021	1996	0	0	1	...	0	0

5 rows × 30 columns


```
In [17]: # Standardize continuous variables

bike_train_df['temp'] = bike_train_df['temp'].apply(lambda t: (t-temp.mean())/temp.std())
bike_train_df['atemp'] = bike_train_df['atemp'].apply(lambda t: (t-atemp.mean())/atemp.std())
bike_train_df['humidity'] = bike_train_df['humidity'].apply(lambda h: (h-humidity.mean())/humidity.std())
bike_train_df['windspeed'] = bike_train_df['windspeed'].apply(lambda w: (w-windspeed.mean())/windspeed.std())

bike_test_df['temp'] = bike_test_df['temp'].apply(lambda t: (t-temp.mean())/temp.std())
bike_test_df['atemp'] = bike_test_df['atemp'].apply(lambda t: (t-atemp.mean())/atemp.std())
bike_test_df['humidity'] = bike_test_df['humidity'].apply(lambda h: (h-humidity.mean())/humidity.std())
bike_test_df['windspeed'] = bike_test_df['windspeed'].apply(lambda w: (w-windspeed.mean())/windspeed.std())

bike_train_df.head()
```

```
Out[17]:
```

	holiday	workingday	temp	atemp	humidity	windspeed	count	season_1	season_2	season_3	...	month_11	day
0	0	1	0.624743	0.651090	0.922058	-0.930164	6073	0	1	0	...	0	0
1	0	1	-0.180583	-0.054841	0.697907	-0.213825	6606	0	0	0	...	0	0
2	0	1	0.803704	0.852785	-0.449062	0.805143	7363	0	1	0	...	0	0
3	0	0	-1.522794	-1.567551	-0.332616	-0.269507	2431	0	0	0	...	0	0
4	0	1	0.535262	0.348548	1.978781	-1.200843	1996	0	0	1	...	0	0

5 rows × 30 columns

1. Why can't they be directly used as predictors?

Because categorical attributes' (those with two or more categories) numbering has no numerical meaning (e.g. "Fall" in season has no connection to the numerical "3"). To use them as predictors, we should only use boolean (0 and 1) to indicate whether such category is presented in the data.

2. Why is it okay to not have a binary column for the K_j -th category?

Because the K_j -th category can be interpreted from the other columns

3. Why shouldn't we include the test set observations in computing the mean and SD?

Because the test set should be "unknown" and its information should not be used in any training purpose

Part (b): Multiple Linear Regression

We are now ready to fit a linear regression model and analyze its coefficients and residuals.

- Fit a multiple linear regression model to the training set, and report its R^2 score on the test set.
- *Statistical significance*: Using a t-test, find out which of estimated coefficients are statistically significant at a significance level of 5% (p-value<0.05). Based on the results of the test, answer the following questions:
 - Which among the predictors have a positive correlation with the number of bike rentals?
 - Does the day of a week have a relationship with bike rentals?
 - Does the month influence the bike rentals?
 - What effect does a holiday have on bike rentals?
 - Is there a difference in the coefficients assigned to temp and atemp? Give an explanation for your observation.
- *Residual plot*: Make a plot of residuals of the fitted model $e = y - \hat{y}$ as a function of the predicted value \hat{y} . Note that this is different from the residual plot for simple linear regression. Draw a horizontal line denoting the zero residual value on the Y-axis. Does the plot reveal a non-linear relationship between the predictors and response? What does the plot convey about the variance of the error terms?

```
In [18]: X_train = sm.add_constant(bike_train_df.drop('count', axis=1).copy())
Y_train = bike_train_df['count'].copy()
```

```
In [19]: MLR_result = OLS(Y_train, X_train).fit()
```

```
In [20]: MLR_result.summary()
```

Out[20]: OLS Regression Results

Dep. Variable:	count	R-squared:	0.576
Model:	OLS	Adj. R-squared:	0.538
Method:	Least Squares	F-statistic:	15.25
Date:	Wed, 04 Oct 2017	Prob (F-statistic):	6.56e-42
Time:	15:42:38	Log-Likelihood:	-2832.1
No. Observations:	331	AIC:	5720.
Df Residuals:	303	BIC:	5827.
Df Model:	27		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	3192.2092	338.813	9.422	0.000	2525.486	3858.933
holiday	-284.3563	398.982	-0.713	0.477	-1069.483	500.770
workingday	308.1531	165.413	1.863	0.063	-17.351	633.657
temp	924.3344	473.819	1.951	0.052	-8.058	1856.727
atemp	311.9618	429.337	0.727	0.468	-532.898	1156.822
humidity	-547.6638	113.029	-4.845	0.000	-770.085	-325.243
windspeed	-254.7369	80.644	-3.159	0.002	-413.431	-96.043
season_1	-1226.1865	506.763	-2.420	0.016	-2223.407	-228.966
season_2	-327.3575	573.373	-0.571	0.568	-1455.654	800.939
season_3	-193.3050	449.171	-0.430	0.667	-1077.194	690.584
month_1	118.8358	505.353	0.235	0.814	-875.611	1113.282
month_2	207.7759	516.216	0.402	0.688	-808.047	1223.599
month_3	358.0167	511.391	0.700	0.484	-648.310	1364.344
month_4	452.1849	657.792	0.687	0.492	-842.234	1746.604
month_5	53.0233	700.991	0.076	0.940	-1326.403	1432.450
month_6	-673.4271	696.142	-0.967	0.334	-2043.313	696.458
month_7	-1161.1512	701.261	-1.656	0.099	-2541.109	218.806
month_8	-657.6397	684.628	-0.961	0.338	-2004.868	689.588
month_9	523.9804	548.284	0.956	0.340	-554.945	1602.906
month_10	605.0867	439.844	1.376	0.170	-260.449	1470.623
month_11	231.5175	413.966	0.559	0.576	-583.094	1046.129
day_of_week_1	-123.7515	170.981	-0.724	0.470	-460.212	212.709
day_of_week_2	-195.2859	203.703	-0.959	0.338	-596.137	205.565
day_of_week_3	170.5113	213.789	0.798	0.426	-250.187	591.210
day_of_week_4	61.2560	206.213	0.297	0.767	-344.536	467.048
day_of_week_5	111.0669	198.877	0.558	0.577	-280.288	502.422
day_of_week_6	465.1450	269.154	1.728	0.085	-64.504	994.794
weather_1	1596.9180	198.428	8.048	0.000	1206.446	1987.390
weather_2	1580.3514	183.091	8.631	0.000	1220.060	1940.643
weather_3	14.9397	386.998	0.039	0.969	-746.605	776.485

Omnibus:	28.947	Durbin-Watson:	1.912
Prob(Omnibus):	0.000	Jarque-Bera (JB):	9.753

Skew:	0.054	Prob(JB):	0.00762
Kurtosis:	2.166	Cond. No.	1.35e+16

```
In [21]: X_test = sm.add_constant(bike_test_df.drop('count', axis=1).copy())
Y_test = bike_test_df['count'].copy()
```

```
In [22]: Y_train_pred = MLR_result.predict(X_train)
r2_train_all = r2_score(Y_train, Y_train_pred)

Y_test_pred = MLR_result.predict(X_test)
r2_test_all = r2_score(Y_test, Y_test_pred)

print("Training R^2: %.5f\nTesting R^2: %.5f" % (r2_train_all, r2_test_all))

Training R^2: 0.57613
Testing R^2: 0.24934
```

```
In [23]: MLR_result.pvalues[MLR_result.pvalues < 0.05]
```

```
Out[23]: const      1.182848e-18
humidity    2.020827e-06
windspeed   1.744235e-03
season_1    1.612408e-02
weather_1    1.930751e-14
weather_2    3.489795e-16
dtype: float64
```

1. Which among the predictors have a positive correlation with the number of bike rentals?

Humidity, Windspeed, Season 1 (Jan), and weathers 1 and 2

2. Does the day of a week have a relationship with bike rentals?

Based on the p-value result - no

3. Does the month influence the bike rentals?

Based on the p-value result - no

4. What effect does a holiday have on bike rentals?

Based on the p-value result - whether it is a holiday has no effect on bike rental

5. Is there a difference in the coefficients assigned to temp and atemp? Give an explanation for your observation.

Yes - the coefficient for temp is roughly 3 times of the atemp. This could be due to the fact that temperature, instead of apparent temperature, is what potential renters use to judge whether it is suitable for biking.

Part (c): Checking Collinearity

Does the data suffer from multi-collinearity? To answer this question, let us first analyze the correlation matrix for the data. Compute the (Pearson product-moment) correlation matrix for the predictor variables in the training set, and visualize the matrix using a heatmap. For categorical attributes, you should use each binary predictor resulting from one-hot encoding to compute their correlations. Are there predictors that fall into natural groups based on the correlation values?


Hint: You may use the `np.corrcoef` function to compute the correlation matrix for a data set (do not forget to transpose the data matrix). You may use `plt.pcolor` function to visualize the correlation matrix.

```
In [24]: corr = bike_train_df.corr()
```

```
In [25]: sns.reset_defaults()
```

```
plt.figure(figsize=(27,18))
plt.xticks(np.arange(len(bike_train_df.columns))+0.5, bike_train_df.columns, rotation=45, \
           fontsize=20, horizontalalignment="right")
plt.yticks(np.arange(len(bike_train_df.columns))+0.5, bike_train_df.columns, fontsize=20)
pc = plt.pcolor(corr, cmap=plt.cm.bwr)
cbar = plt.colorbar(pc)
cbar.ax.set_yticklabels(labels=cbar.ax.get_yticklabels(), fontdict={"size": 20})

plt.show()
```



Some of the predictors are highly correlated to each others:

For example, workingday and holiday are obviously related to each other. Seasons, months and temperature are related to each others. Weather and humidity are another pair with high correlation

So these predictors do not fall into natural group as some of them are highly correlated

Part (d): Subset Selection

Apply either one of the following subset selection methods discussed in class to choose a minimal subset of predictors that are related to the response variable:

- Step-wise forward selection
- Step-wise backward selection

We require you to implement both these methods *from scratch*. You may use the Bayesian Information Criterion (BIC) to choose the subset size in each method. Do these methods eliminate one or more of the redundant predictors (if any) identified in Part (c)? In each case, fit linear regression models using the identified subset of predictors to the training set. How do the test R^2 scores for the fitted models compare with the model fitted in Part (b) using all predictors?

```
In [26]: all_predictors = bike_train_df.drop('count', axis=1).columns.tolist()
```

Forward Selection

```
In [27]: predictors = ([[], 0]) # (predictors, bic)
Y_train = bike_train_df['count']

for k in range(1, len(all_predictors)+1):
    best_k_minus_1 = predictors[-1][0]
    # Get the list of remaining predictors
    new_predictors = list(set(all_predictors) - set(best_k_minus_1))
    bics = []

    for predictor in new_predictors:
        k_predictors = best_k_minus_1 + [predictor]
        k_X_train = sm.add_constant(bike_train_df[k_predictors])

        k_MLR_result = OLS(Y_train, k_X_train).fit()
        bics.append(k_MLR_result.bic)

    best_k = best_k_minus_1 + [new_predictors[np.argmin(bics)]]
    predictors.append((best_k, np.min(bics)))
```

```
In [28]: best_predictors_forward = sorted(predictors, key=lambda p: p[1])[1]
print("The best predictors (forward selection) set is:\n%s\nwith BIC=%.3f"
      % (best_predictors_forward[0], best_predictors_forward[1]))
```

The best predictors (forward selection) set is:
['atemp', 'humidity', 'season_1', 'weather_3', 'month_9', 'month_10', 'windspeed', 'month_7']
with BIC=5737.502

```
In [29]: best_X_train_forward = sm.add_constant(bike_train_df[best_predictors_forward[0]])

best_MLR_model = OLS(Y_train, best_X_train_forward)
best_MLR_result = best_MLR_model.fit()

Y_train_pred = best_MLR_result.predict(best_X_train_forward)
r2_train_forward = r2_score(Y_train, Y_train_pred)

best_X_test_forward = sm.add_constant(bike_test_df[best_predictors_forward[0]])
Y_test = bike_test_df['count']

Y_test_pred = best_MLR_result.predict(best_X_test_forward)
r2_test_forward = r2_score(Y_test, Y_test_pred)

print("R^2 score for training set: %.3f\nR^2 score for testing set: %.3f" % (r2_train_forward,
r2_test_forward))

R^2 score for training set: 0.548
R^2 score for testing set: 0.266
```

Backward Selection

```

In [30]: all_X_train = sm.add_constant(bike_train_df[all_predictors])
Y_train = bike_train_df['count']
all_MLR_result = OLS(Y_train, all_X_train).fit()
all_bic = all_MLR_result.bic

predictors = [(all_predictors, all_bic)] # (predictors, bic)

for k in range(len(all_predictors), 1, -1):
    best_k = predictors[-1][0]
    bics = []

    for predictor in best_k:
        # Drop one of the predictors and test the bic
        k_minus_1 = list(set(best_k) - set([predictor]))
        k_X_train = sm.add_constant(bike_train_df[k_minus_1])

        k_MLR_result = OLS(Y_train, k_X_train).fit()
        bics.append(k_MLR_result.bic)

    best_k_minus_1 = list(set(best_k) - set([best_k[np.argmin(bics)]]))
    predictors.append((best_k_minus_1, np.min(bics)))

```

```

In [31]: best_predictors_backward = sorted(predictors, key=lambda p: p[1])[0]
print("The best predictors set (backward selection) is:\n%s\nwith BIC=%.3f"
      % (best_predictors_backward[0], best_predictors_backward[1]))

```

The best predictors set (backward selection) is:
['weather_3', 'season_1', 'humidity', 'month_8', 'month_6', 'temp', 'windspeed', 'month_7']
with BIC=5736.208

```

In [32]: best_X_train_backward = sm.add_constant(bike_train_df[best_predictors_backward[0]])

best_MLR_model = OLS(Y_train, best_X_train_backward)
best_MLR_result = best_MLR_model.fit()

Y_train_pred = best_MLR_result.predict(best_X_train_backward)
r2_train_backward = r2_score(Y_train, Y_train_pred)

best_X_test_backward = sm.add_constant(bike_test_df[best_predictors_backward[0]])
Y_test = bike_test_df['count']

Y_test_pred = best_MLR_result.predict(best_X_test_backward)
r2_test_backward = r2_score(Y_test, Y_test_pred)

print("R^2 score for training set: %.3f\nR^2 score for testing set: %.3f" % (r2_train_backward, r2_test_backward))

R^2 score for training set: 0.550
R^2 score for testing set: 0.264

```

```

In [33]: print("All predictors included: R^2 training: %.3f, R^2 testing: %.3f\n" % (r2_train_all, r2_test_all))
print("Best forward selection set (R^2 training: %.3f, R^2 testing: %.3f):\n\t%s\n"
      % (r2_train_forward, r2_test_forward, best_predictors_forward))
print("Best backward selection set (R^2 training: %.3f, R^2 testing: %.3f):\n\t%s\n"
      % (r2_train_backward, r2_test_backward, best_predictors_backward))

All predictors included: R^2 training: 0.576, R^2 testing: 0.249

Best forward selection set (R^2 training: 0.548, R^2 testing: 0.266):
(['atemp', 'humidity', 'season_1', 'weather_3', 'month_9', 'month_10', 'windspeed', 'month_7'], 5737.5016073504503)

Best backward selection set (R^2 training: 0.550, R^2 testing: 0.264):
(['weather_3', 'season_1', 'humidity', 'month_8', 'month_6', 'temp', 'windspeed', 'month_7'], 5736.2082447442772)

```

Both stepwise selection (forward and backward) eliminate some of the redundant (correlated) predictors found in part c), reducing the BIC score. The R^2 scores for the training set decrease in both cases, which is expected due to less predictors included. The R^2 scores for the testing set increase in both cases, which is also expected since overfitting problem is resolved.

Part (e): Cross Validation

- Perform a 10-fold cross-validation procedure to select between the 3 competing models you have so far: the model with the best BIC from Step-wise forward selection, the model with the best BIC from Step-wise backward selection (if it is different), and the model with all possible predictors. Report the average R^2 across all 10 validation sets for each model and compare the results. Why do you think this is the case?
- Evaluate each of the 3 models on the provided left out test set by calculating R^2 . Do the results agree with the cross-validation? Why or why not?

```
In [34]: from sklearn.model_selection import KFold
```

```
In [35]: n_folds = 10
predictors_set = [all_predictors, best_predictors_forward[0], best_predictors_backward[0]]
r2_valid = []

for train_ind, valid_ind in KFold(n_folds, shuffle=True).split(bike_train_df):
    # Split the training set into 10 portion - 9 for training and 1 for validation
    Y_train_cv = bike_train_df.loc[train_ind, 'count']
    Y_valid_cv = bike_train_df.loc[valid_ind, 'count']

    for predictors in predictors_set:
        X_train_cv = sm.add_constant(bike_train_df.loc[train_ind, predictors])
        X_valid_cv = sm.add_constant(bike_train_df.loc[valid_ind, predictors])

        MLR_result = OLS(Y_train_cv, X_train_cv).fit()
        Y_valid_pred = MLR_result.predict(X_valid_cv)
        r2_valid.append(r2_score(Y_valid_cv, Y_valid_pred))

r2_valid = np.array(r2_valid).reshape((n_folds, 3))
r2_valid_mean = r2_valid.mean(axis=0)
```

```
In [36]: print("Average Cross-Validation R^2 score for all predictors: %.3f" % r2_valid_mean[0])
print("Average Cross-Validation R^2 score for forward selection: %.3f" % r2_valid_mean[1])
print("Average Cross-Validation R^2 score for backward selection: %.3f" % r2_valid_mean[2])
```

```
Average Cross-Validation R^2 score for all predictors: 0.389
Average Cross-Validation R^2 score for forward selection: 0.498
Average Cross-Validation R^2 score for backward selection: 0.497
```

The validation R^2 score using all predictors is lower simply because of overfitting. R^2 scores for forward and backward selection are similar, which is expected as they eliminate most correlated predictors and reduce the order of overfitting.

```
In [37]: r2_test = []
Y_train = bike_train_df['count']
Y_test = bike_test_df['count']

for predictors in predictors_set:
    X_train = sm.add_constant(bike_train_df[predictors])
    X_test = sm.add_constant(bike_test_df[predictors])

    MLR_result = OLS(Y_train, X_train).fit()
    Y_test_pred = MLR_result.predict(X_test)
    r2_test.append(r2_score(Y_test, Y_test_pred))
```

```
In [38]: print("Testing Set R^2 score for all predictors: %.3f" % r2_test[0])
print("Testing Set R^2 score for forward selection: %.3f" % r2_test[1])
print("Testing Set R^2 score for backward selection: %.3f" % r2_test[2])
```

```
Testing Set R^2 score for all predictors: 0.249
Testing Set R^2 score for forward selection: 0.266
Testing Set R^2 score for backward selection: 0.264
```

The testing results agree with the cross-validation result, where the R^2 scores for both forward and backward selection are similar and higher (i.e. better fitting) than that fitted from all predictors

```
In [ ]:
```